

Microservice Rate Limiting System Design	
Functional Requirements Limit each client (by IP or token) to a fixed number of requests per time window Block excessive requests with an HTTP 429 response Apply rate limiting globally or per endpoint as needed Log and monitor rejected requests	Non-Functional Requirements Minimal latency overhead (should not slow down response time) Lightweight for single-service deployment (no extra infra required) Optional support for distributed limit coordination (e.g., Redis) Easy to configure and maintain Focus is on Consistency and Availability . It's important that request counts are accurate and consistently enforced across all replicas — otherwise abuse may slip through. Partition Tolerance is less critical, since this service runs best within a single region or tightly connected cluster.
Chosen Algorithm: Token Bucket Allows bursts while enforcing average request rates Replenishes tokens at a fixed rate over time Denies request when bucket is empty (returns 429)	
Architecture Overview	
Clients	Any API consumer (browser, mobile app, back
Entry Layer	FastAPI App with endpoint /api/v1/data All requests pass through rate limiter
Rate Limiter Middleware	Integrated via slowapi Keeps token count per client/IP Refill rate: 1 token/second Max tokens: 100 (100 requests per 100 sec)
Application Logic	If allowed: returns JSON success response If limited: returns HTTP 429 with error messag
Optional	Redis (if scaling across multiple instances) Prometheus/Grafana for monitoring
Flow Diagram (Text Format) Client → FastAPI App → RateLimiter Middleware → <div><div>└─ Enough tokens → Handle request → 200 OK</div><div>└─ No tokens left → Return 429 Too Many Requests</div></div>	
Link to GitHub: https://github.com/jaxylykovm/rate-limit-microservice	

